

# Chips Sales Analysis

- Created by: Rupsa Chaudhuri
- LinkedIn: <https://www.linkedin.com/in/rupsa-chaudhuri/>
- GitHub: <https://github.com/rupsa723?tab=repositories>

## Load the Data

We'll read the data from the provided files:

QVI\_transaction\_data.xlsx: This file likely contains transactional data.

QVI\_purchase\_behaviour.csv: This file likely contains information on purchase behavior.

```
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from datetime import datetime
import re

# Load the data
transaction_data = pd.read_excel(r'C:\Users\rupsa\Downloads\Quantium Internship\QVI_transaction_data.xlsx')
purchase_behaviour = pd.read_csv(r'C:\Users\rupsa\Downloads\Quantium Internship\QVI_purchase_behaviour.csv')

# Display the first few rows of each dataframe
print("Transaction Data:")
print(transaction_data.head())

print("\nPurchase Behaviour Data:")
print(purchase_behaviour.head())
```

Transaction Data:

	DATE	STORE_NBR	LYLTY_CARD_NBR	TXN_ID	PROD_NBR	\
0	43390	1	1000	1	5	
1	43599	1	1307	348	66	
2	43605	1	1343	383	61	
3	43329	2	2373	974	69	
4	43330	2	2426	1038	108	

	PROD_NAME	PROD_QTY	TOT_SALES
0	Natural Chip Compny SeaSalt175g	2	6.0
1	CCs Nacho Cheese 175g	3	6.3
2	Smiths Crinkle Cut Chips Chicken 170g	2	2.9
3	Smiths Chip Thinly S/Cream&Onion 175g	5	15.0
4	Kettle Tortilla ChpsHny&Jlpno Chili 150g	3	13.8

## Purchase Behaviour Data:

	LYLTY_CARD_NBR		LIFESTAGE	PREMIUM_CUSTOMER
0	1000	YOUNG	SINGLES/COUPLES	Premium
1	1002	YOUNG	SINGLES/COUPLES	Mainstream
2	1003		YOUNG FAMILIES	Budget
3	1004	OLDER	SINGLES/COUPLES	Mainstream
4	1005	MIDAGE	SINGLES/COUPLES	Mainstream

## Summarize the Data

We'll generate basic summaries to get an initial understanding of the data.

```
# Summary statistics for transaction data
```

```
transaction_summary = transaction_data.describe(include='all')
```

```
# Summary statistics for purchase behaviour data
```

```
purchase_behaviour_summary =
```

```
purchase_behaviour.describe(include='all')
```

```
transaction_summary, purchase_behaviour_summary
```

	DATE	STORE_NBR	LYLTY_CARD_NBR	TXN_ID	\
count	264836.000000	264836.000000	2.648360e+05	2.648360e+05	
unique	NaN	NaN	NaN	NaN	
top	NaN	NaN	NaN	NaN	
freq	NaN	NaN	NaN	NaN	
mean	43464.036260	135.08011	1.355495e+05	1.351583e+05	
std	105.389282	76.78418	8.057998e+04	7.813303e+04	
min	43282.000000	1.00000	1.000000e+03	1.000000e+00	
25%	43373.000000	70.00000	7.002100e+04	6.760150e+04	
50%	43464.000000	130.00000	1.303575e+05	1.351375e+05	
75%	43555.000000	203.00000	2.030942e+05	2.027012e+05	
max	43646.000000	272.00000	2.373711e+06	2.415841e+06	

	PROD_NBR	PROD_NAME
PROD_QTY \		
count	264836.000000	264836
unique	NaN	114
top	NaN	Kettle Mozzarella Basil & Pesto 175g
freq	NaN	3304
mean	56.583157	NaN
std	32.826638	NaN
min	1.000000	NaN

1.000000		
25%	28.000000	NaN
2.000000		
50%	56.000000	NaN
2.000000		
75%	85.000000	NaN
2.000000		
max	114.000000	NaN
200.000000		

	TOT_SALES
count	264836.000000
unique	NaN
top	NaN
freq	NaN
mean	7.304200
std	3.083226
min	1.500000
25%	5.400000
50%	7.400000
75%	9.200000
max	650.000000

	LYLTY_CARD_NBR	LIFESTAGE	PREMIUM_CUSTOMER
count	7.263700e+04	72637	72637
unique	NaN	7	3
top	NaN	RETIREEES	Mainstream
freq	NaN	14805	29245
mean	1.361859e+05	NaN	NaN
std	8.989293e+04	NaN	NaN
min	1.000000e+03	NaN	NaN
25%	6.620200e+04	NaN	NaN
50%	1.340400e+05	NaN	NaN
75%	2.033750e+05	NaN	NaN
max	2.373711e+06	NaN	NaN

## Identify Missing Values

```
# Check for missing values in transaction data
missing_values_transaction = transaction_data.isnull().sum()

# Check for missing values in purchase behaviour data
missing_values_behaviour = purchase_behaviour.isnull().sum()

missing_values_transaction, missing_values_behaviour

(DATE          0
STORE_NBR      0
LYLTY_CARD_NBR 0
TXN_ID         0
PROD_NBR       0)
```


```
PROD_NAME      0
PROD_QTY       0
TOT_SALES      0
dtype: int64,
LYLTY_CARD_NBR 0
LIFESTAGE      0
PREMIUM_CUSTOMER 0
dtype: int64)
```

## Identifying and Handling Outliers

TOT\_SALES: Mostly small transactions, with an average of 7.30 and a maximum of 650.00 dollars, indicating potential outliers. We'll focus on TOT\_SALES to identify potential outliers.

```
import matplotlib.pyplot as plt

# Plot TOT_SALES to identify outliers
plt.figure(figsize=(6,4))
plt.boxplot(transaction_data['TOT_SALES'])
plt.title('Boxplot of TOT_SALES')
plt.ylabel('Total Sales')
plt.show()
```



The boxplot of TOT\_SALES indicates that there are indeed some outliers, with most transactions clustered below 10 and a few extending significantly higher, up to 650.00.

We will remove the outliers by filtering out top 1% of the data.

```
# Calculate the 99th percentile
threshold = transaction_data['TOT_SALES'].quantile(0.99)
# Filter the data to remove the top 1%
filtered_transaction_data =
transaction_data[transaction_data['TOT_SALES'] <= threshold]

print("\nFiltered Data:")
print(filtered_transaction_data.describe())
```

Filtered Data:

	DATE	STORE_NBR	LYLTY_CARD_NBR	TXN_ID \
count	264051.000000	264051.000000	2.640510e+05	2.640510e+05
mean	43464.046074	135.067714	1.355380e+05	1.351460e+05
std	105.282999	76.786795	8.059354e+04	7.813593e+04
min	43282.000000	1.000000	1.000000e+03	1.000000e+00
25%	43373.000000	70.000000	7.001800e+04	6.758450e+04
50%	43464.000000	130.000000	1.303540e+05	1.351170e+05
75%	43555.000000	203.000000	2.030930e+05	2.026935e+05
max	43646.000000	272.000000	2.373711e+06	2.415841e+06

	PROD_NBR	PROD_QTY	TOT_SALES
count	264051.000000	264051.000000	264051.000000
mean	56.593170	1.898796	7.267126
std	32.824607	0.316009	2.450433
min	1.000000	1.000000	1.500000
25%	28.000000	2.000000	5.400000
50%	56.000000	2.000000	7.400000
75%	85.000000	2.000000	8.800000
max	114.000000	5.000000	13.000000

## Check and Correct Data Formats

```
# Convert DATE to datetime format
filtered_transaction_data['DATE'] =
pd.to_datetime(filtered_transaction_data['DATE'], origin='1899-12-30',
unit='D')
```

```
# Verify the date format
print(filtered_transaction_data['DATE'].head())
```

```
0    2018-10-17
1    2019-05-14
2    2019-05-20
5    2019-05-19
6    2019-05-16
```

```
Name: DATE, dtype: datetime64[ns]
```

```
C:\Users\rupsa\AppData\Local\Temp\ipykernel_19920\1941825135.py:2:
SettingWithCopyWarning:
```

A value is trying to be set on a copy of a slice from a DataFrame.  
Try using `.loc[row_indexer,col_indexer] = value` instead

See the caveats in the documentation:

[https://pandas.pydata.org/pandas-docs/stable/user\\_guide/indexing.html#returning-a-view-versus-a-copy](https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy)

```
filtered_transaction_data['DATE'] =  
pd.to_datetime(filtered_transaction_data['DATE'], origin='1899-12-30',  
unit='D')
```

## Derive Additional Features

Extract pack size and brand name:

```
# Extract pack size from PROD_NAME  
filtered_transaction_data['PACK_SIZE'] =  
filtered_transaction_data['PROD_NAME'].str.extract(r'(\d+)g')  
filtered_transaction_data['PACK_SIZE'] =  
filtered_transaction_data['PACK_SIZE'].fillna(0).astype(int)
```

```
# Extract brand name
```

```
filtered_transaction_data['BRAND_NAME'] =  
filtered_transaction_data['PROD_NAME'].str.split().str[0]
```

C:\Users\rupsa\AppData\Local\Temp\ipykernel\_19920\1811567043.py:2:  
SettingWithCopyWarning:

A value is trying to be set on a copy of a slice from a DataFrame.  
Try using `.loc[row_indexer,col_indexer] = value` instead

See the caveats in the documentation:

[https://pandas.pydata.org/pandas-docs/stable/user\\_guide/indexing.html#returning-a-view-versus-a-copy](https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy)

```
filtered_transaction_data['PACK_SIZE'] =  
filtered_transaction_data['PROD_NAME'].str.extract(r'(\d+)g')
```

C:\Users\rupsa\AppData\Local\Temp\ipykernel\_19920\1811567043.py:3:  
SettingWithCopyWarning:

A value is trying to be set on a copy of a slice from a DataFrame.  
Try using `.loc[row_indexer,col_indexer] = value` instead

See the caveats in the documentation:

[https://pandas.pydata.org/pandas-docs/stable/user\\_guide/indexing.html#returning-a-view-versus-a-copy](https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy)

```
filtered_transaction_data['PACK_SIZE'] =  
filtered_transaction_data['PACK_SIZE'].fillna(0).astype(int)
```

C:\Users\rupsa\AppData\Local\Temp\ipykernel\_19920\1811567043.py:6:  
SettingWithCopyWarning:

A value is trying to be set on a copy of a slice from a DataFrame.  
Try using `.loc[row_indexer,col_indexer] = value` instead

See the caveats in the documentation:

```
https://pandas.pydata.org/pandas-docs/stable/user\_guide/indexing.html#returning-a-view-versus-a-copy
filtered_transaction_data['BRAND_NAME'] =
filtered_transaction_data['PROD_NAME'].str.split().str[0]
```

```
# Check the results
```

```
brand_summary = filtered_transaction_data['BRAND_NAME'].value_counts()
brand_summary
```

```
BRAND_NAME
Kettle      41108
Smiths      28791
Pringles    25012
Doritos     24875
Thins       14029
RRD         11874
Infuzions   11035
WW          10320
Cobs        9669
Tostitos    9417
Twisties    9411
Old         9284
Tyrrells    6428
Grain       6258
Natural     6037
Red         5870
Cheezels    4583
CCs         4551
Woolworths  4437
Dorito      3175
Infzns      3138
Smith       2963
Cheetos     2915
Snbts       1576
Burger      1564
GrnWves     1465
Sunbites    1432
French      1418
NCC         1416
Name: count, dtype: int64
```

Clean brand names

```
# Convert brand names to a consistent case (e.g., uppercase)
```

```
filtered_transaction_data['BRAND_NAME'] =
filtered_transaction_data['BRAND_NAME'].str.upper()
```

```
# Comprehensive replacement dictionary
```

```
brand_replacements = {
```

```

    'RED': 'RRD',
    'SNBTS': 'SUNBITES',
    'INFZNS': 'INFUZIONI',
    'WW': 'WOOLWORTHS',
    'SMITH': 'SMITHS',
    'NCC': 'NATURAL',
    'DORITO': 'DORITOS',
    'GRAIN': 'GRNWVES',
    'WOOLWORTHS': 'WOOLWORTHS' # Ensure consistency
}

# Replace brand names using the dictionary
filtered_transaction_data['BRAND_NAME'] =
filtered_transaction_data['BRAND_NAME'].replace(brand_replacements)

# Display value counts of the cleaned brand names
print(filtered_transaction_data['BRAND_NAME'].value_counts())

```

```

BRAND_NAME
KETTLE      41108
SMITHS      31754
DORITOS     28050
PRINGLES    25012
RRD         17744
WOOLWORTHS  14757
INFUZIONI   14173
THINS       14029
COBS        9669
TOSTITOS    9417
TWISTIES    9411
OLD         9284
GRNWVES     7723
NATURAL     7453
TYRRELLS    6428
CHEEZELS    4583
CCS         4551
SUNBITES    3008
CHEETOS     2915
BURGER      1564
FRENCH      1418
Name: count, dtype: int64

```

```

# Verify new features
print(filtered_transaction_data[['PROD_NAME', 'PACK_SIZE',
'BRAND_NAME']].head())


```

		PROD_NAME	PACK_SIZE	BRAND_NAME
0	Natural Chip	Compny SeaSalt175g	175	NATURAL
1		CCs Nacho Cheese 175g	175	CCS
2	Smiths Crinkle Cut	Chips Chicken 170g	170	SMITHS

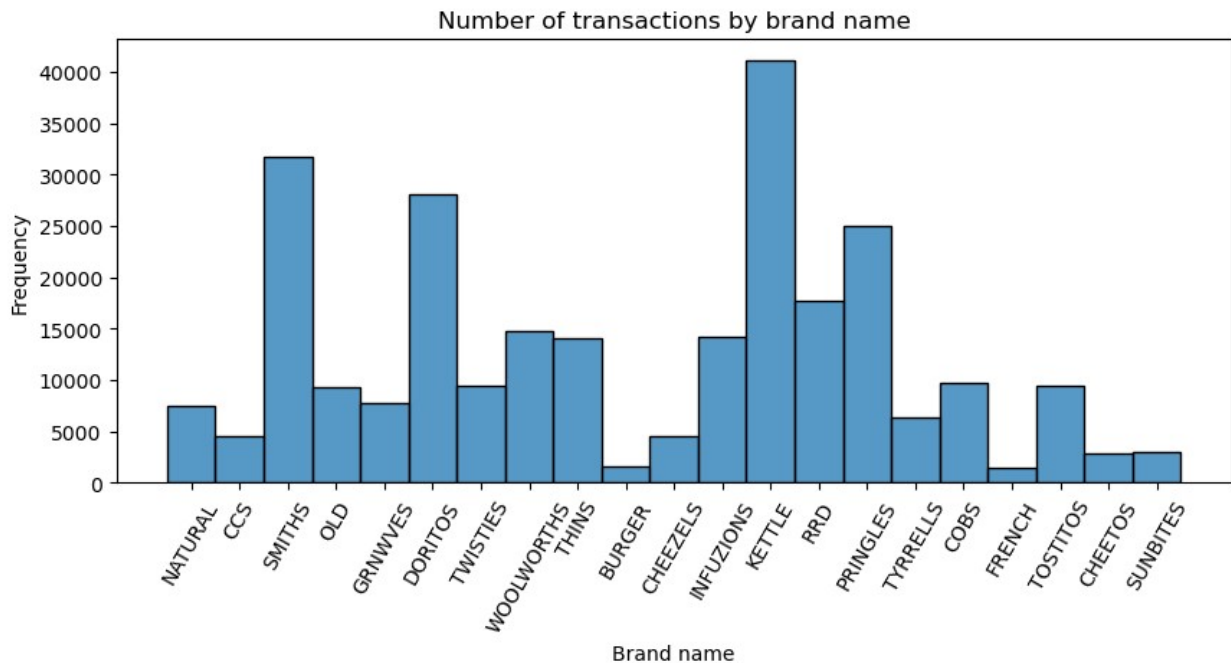


5	Old El Paso Salsa	Dip Tomato Mild	300g	300	OLD
6	Smiths Crinkle Chips	Salt & Vinegar	330g	330	SMITHS

```
# Plot a histogram of PACK_SIZE
plt.figure(figsize=(6,4))
sns.histplot(filtered_transaction_data['PACK_SIZE'], bins=range(50,
400, 10))
plt.title('Number of transactions by pack size')
plt.xlabel('Pack Size (g)')
plt.ylabel('Frequency')
plt.show()
```



```
# Plot a histogram of BRAND_NAME
plt.figure(figsize=(10,4))
sns.histplot(filtered_transaction_data['BRAND_NAME'], bins=range(50,
400, 10))
plt.title('Number of transactions by brand name')
plt.xlabel('Brand name')
plt.ylabel('Frequency')
plt.xticks(rotation=60)
plt.show()
```



Count the frequency of each word occuring in product name:

```
# Further examine PROD_NAME
product_words =
pd.DataFrame(filtered_transaction_data['PROD_NAME'].str.split(expand=True).stack().value_counts().reset_index())
product_words.columns = ['words', 'frequency']
```

```
# Remove digits and special characters
product_words = product_words[~product_words['words'].str.contains(r'\d|&|\$|%', regex=True)]
print(product_words)
```

	words	frequency
1	Chips	49634
3	Kettle	41108
5	Smiths	28791
6	Salt	27912
7	Cheese	27794
...	...	...
214	Pc	1429
216	Aioli	1427
217	Fries	1418
218	Garden	1416
219	NCC	1416

[171 rows x 2 columns]

## Count the number of transactions by date

```
transactions_by_day =  
filtered_transaction_data.groupby('DATE').size().reset_index(name='N')  
print(transactions_by_day)
```

	DATE	N
0	2018-07-01	724
1	2018-07-02	711
2	2018-07-03	722
3	2018-07-04	714
4	2018-07-05	712
...	...	...
359	2019-06-26	723
360	2019-06-27	709
361	2019-06-28	730
362	2019-06-29	745
363	2019-06-30	744

[364 rows x 2 columns]

```
# Create a sequence of dates and join this with the count of  
transactions by date
```

```
full_date_range = pd.DataFrame({'DATE': pd.date_range(start='2018-07-  
01', end='2019-06-30')})  
transactions_by_day = pd.merge(full_date_range, transactions_by_day,  
on='DATE', how='left').fillna(0)  
print(transactions_by_day)
```


	DATE	N
0	2018-07-01	724.0
1	2018-07-02	711.0
2	2018-07-03	722.0
3	2018-07-04	714.0
4	2018-07-05	712.0
...	...	...
360	2019-06-26	723.0
361	2019-06-27	709.0
362	2019-06-28	730.0
363	2019-06-29	745.0
364	2019-06-30	744.0

[365 rows x 2 columns]

```
# Plot transactions over time
```

```
plt.figure(figsize=(10,4))  
sns.lineplot(data=transactions_by_day, x='DATE', y='N')  
plt.title('Transactions over time')  
plt.xlabel('Day')  
plt.ylabel('Number of transactions')
```

```
plt.xticks(rotation=90)
plt.show()
```




We can see that there is an increase in purchases in December and a break in late December. Let's zoom in on this.

```
# Filter to December
transactions_by_day['DATE'] =
pd.to_datetime(transactions_by_day['DATE'])
december_data =
transactions_by_day[transactions_by_day['DATE'].dt.month == 12]

# Plotting
plt.figure(figsize=(10, 4))
sns.lineplot(data=december_data, x='DATE', y='N')
plt.xlabel('Day')
plt.ylabel('Number of transactions')
plt.title('Transactions over time')
plt.xticks(rotation=45)
plt.gca().xaxis.set_major_locator(plt.MaxNLocator(nbins=len(december_data), prune='both'))

# Show plot
plt.tight_layout()
plt.show()
```



We can see that the increase in sales occurs in the lead-up to Christmas and that there are zero sales on Christmas day itself. This is due to shops being closed on Christmas day.

```
filtered_transaction_data.head()
```

	DATE	STORE_NBR	LYLTY_CARD_NBR	TXN_ID	PROD_NBR	\
0	2018-10-17	1	1000	1	5	
1	2019-05-14	1	1307	348	66	
2	2019-05-20	1	1343	383	61	
5	2019-05-19	4	4074	2982	57	
6	2019-05-16	4	4149	3333	16	

	PROD_NAME	PROD_QTY	TOT_SALES
0	Natural Chip Compny SeaSalt175g	2	6.0
1	CCs Nacho Cheese 175g	3	6.3
2	Smiths Crinkle Cut Chips Chicken 170g	2	2.9
5	Old El Paso Salsa Dip Tomato Mild 300g	1	5.1
6	Smiths Crinkle Chips Salt & Vinegar 330g	1	5.7

	BRAND_NAME
0	NATURAL
1	CCS
2	SMITHS
5	OLD
6	SMITHS

## Correlation Analysis

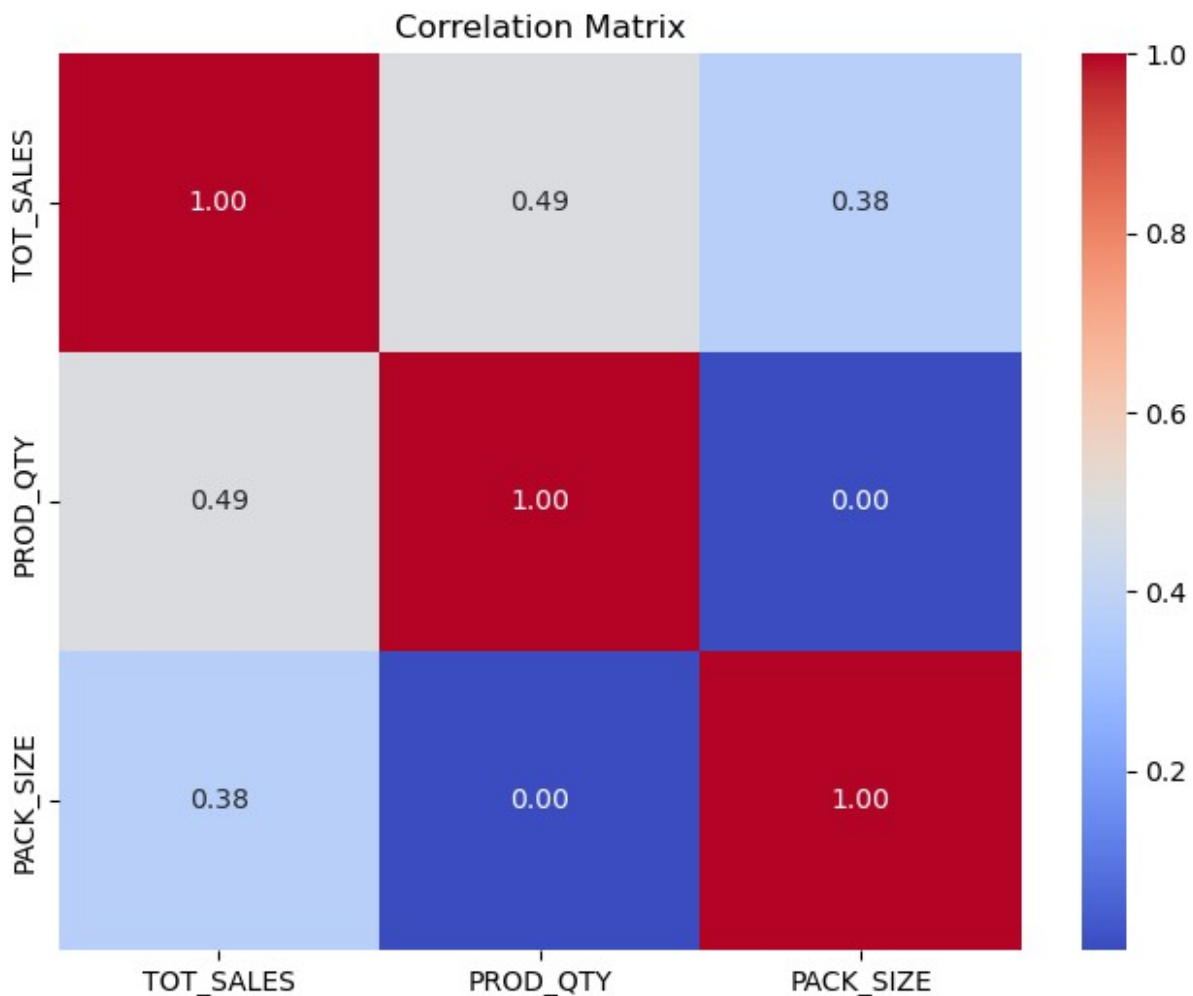
```
# Compute the correlation matrix
corr_matrix = filtered_transaction_data[['TOT_SALES', 'PROD_QTY',
'PACK_SIZE']].corr()
```

```
corr_matrix
```

	TOT_SALES	PROD_QTY	PACK_SIZE
TOT_SALES	1.000000	0.488852	0.375013
PROD_QTY	0.488852	1.000000	0.000074
PACK_SIZE	0.375013	0.000074	1.000000

```
# Plot the correlation matrix
```

```
plt.figure(figsize=(8, 6))
sns.heatmap(corr_matrix, annot=True, cmap='coolwarm', fmt='.2f')
plt.title('Correlation Matrix')
plt.show()
```



- Moderate Positive Correlation: Total sales and product quantity have a moderate positive correlation (0.488).
- Weak Positive Correlation: Total sales and pack size show a weak positive correlation (0.375).
- No Correlation: Product quantity and pack size have no significant correlation (0.000074).

## Customer segmentation

Let's uncover:

Chip spending by customer group (lifestages & budget habits) Number of customers in each lifestage group Average chip purchases per customer group Average chip price per customer group

```
# Merge transaction and purchase behaviour data
merged_data = filtered_transaction_data.merge(purchase_behaviour,
on='LYLTY_CARD_NBR',how='left')

# Calculate total sales by LIFESTAGE and PREMIUM_CUSTOMER
total_sales_by_segment = merged_data.groupby(['LIFESTAGE',
'PREMIUM_CUSTOMER'])['TOT_SALES'].sum().reset_index()

# Verify the results
print(total_sales_by_segment)
```


		LIFESTAGE	PREMIUM_CUSTOMER	TOT_SALES
0	MIDAGE	SINGLES/COUPLES	Budget	35239.70
1	MIDAGE	SINGLES/COUPLES	Mainstream	90082.00
2	MIDAGE	SINGLES/COUPLES	Premium	58041.45
3		NEW FAMILIES	Budget	21847.85
4		NEW FAMILIES	Mainstream	16927.15
5		NEW FAMILIES	Premium	11450.50
6		OLDER FAMILIES	Budget	166966.45
7		OLDER FAMILIES	Mainstream	102457.90
8		OLDER FAMILIES	Premium	79937.20
9	OLDER	SINGLES/COUPLES	Budget	135676.80
10	OLDER	SINGLES/COUPLES	Mainstream	132407.00
11	OLDER	SINGLES/COUPLES	Premium	131260.25
12		RETIREEES	Budget	112263.60
13		RETIREEES	Mainstream	154774.35
14		RETIREEES	Premium	96926.40
15		YOUNG FAMILIES	Budget	138196.20
16		YOUNG FAMILIES	Mainstream	92000.35
17		YOUNG FAMILIES	Premium	83428.40
18	YOUNG	SINGLES/COUPLES	Budget	60840.30
19	YOUNG	SINGLES/COUPLES	Mainstream	156716.50
20	YOUNG	SINGLES/COUPLES	Premium	41451.60

```
# Check for missing customer details
missing_customers = merged_data[merged_data.isnull().any(axis=1)]
missing_customers

Empty DataFrame
Columns: [DATE, STORE_NBR, LYLTY_CARD_NBR, TXN_ID, PROD_NBR,
PROD_NAME, PROD_QTY, TOT_SALES, PACK_SIZE, BRAND_NAME, LIFESTAGE,
PREMIUM_CUSTOMER]
Index: []

# Plot total sales by LIFESTAGE and PREMIUM_CUSTOMER
plt.figure(figsize=(6, 4))
sns.barplot(data=total_sales_by_segment, x='LIFESTAGE', y='TOT_SALES',
hue='PREMIUM_CUSTOMER')
plt.title('Total sales by LIFESTAGE and PREMIUM_CUSTOMER')
plt.xticks(rotation=70)
plt.ylabel('Total Sales')
plt.show()
```



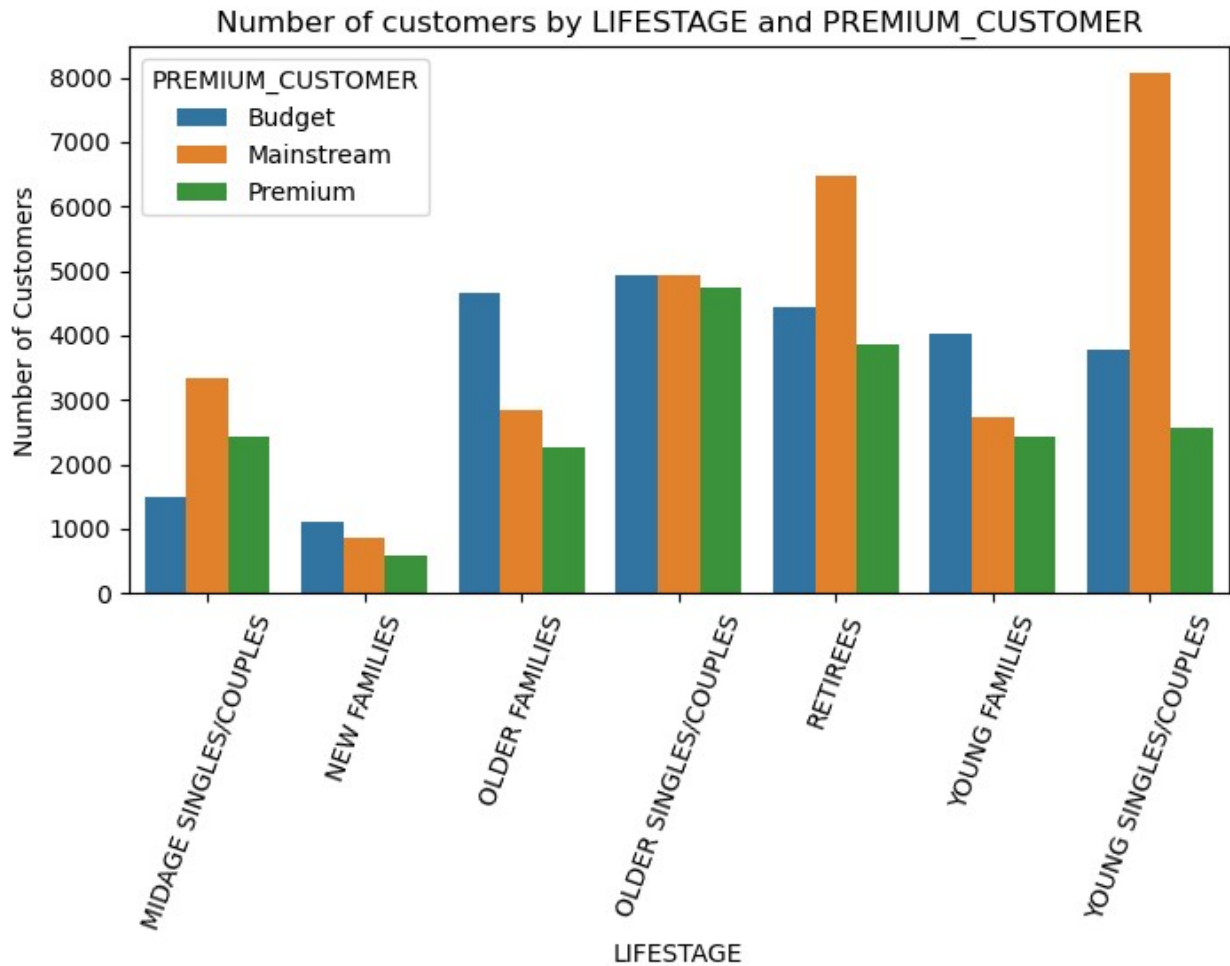


Sales are coming mainly from Budget - older families, Mainstream - young singles/couples, and Mainstream- retirees

Let's see if the higher sales are due to there being more customers who buy chips

```
# Number of customers by LIFESTAGE and PREMIUM_CUSTOMER
customers_by_segment = merged_data.groupby(['LIFESTAGE',
'PREMIUM_CUSTOMER'])['LYLTY_CARD_NBR'].nunique().reset_index()

# Plot number of customers by LIFESTAGE and PREMIUM_CUSTOMER
plt.figure(figsize=(8, 4))
sns.barplot(data=customers_by_segment, x='LIFESTAGE',
y='LYLTY_CARD_NBR', hue='PREMIUM_CUSTOMER')
plt.title('Number of customers by LIFESTAGE and PREMIUM_CUSTOMER')
plt.xticks(rotation=70)
plt.ylabel('Number of Customers')
plt.show()
```



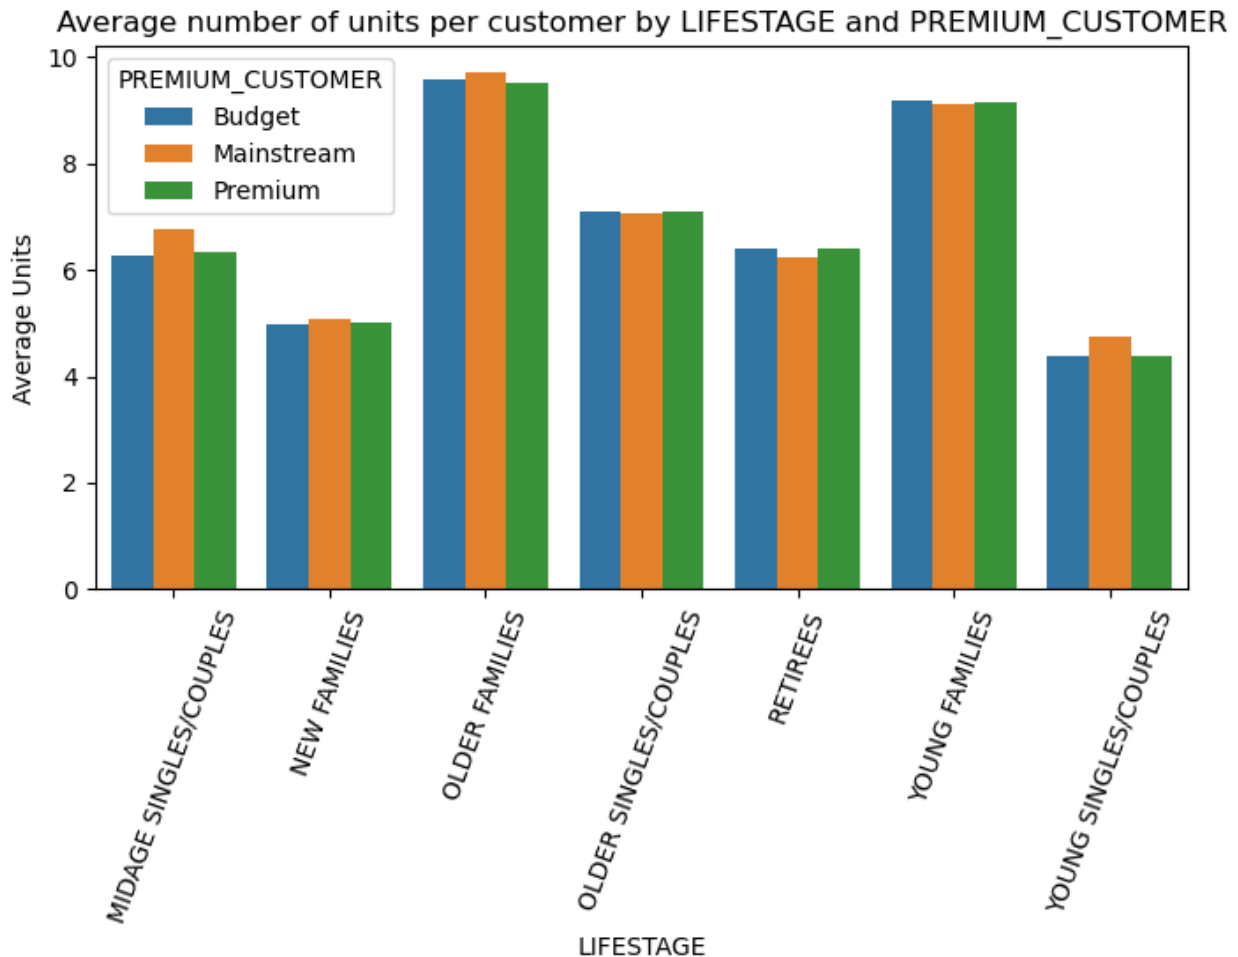
There are more Mainstream - young singles/couples and Mainstream - retirees who buy chips. This contributes to there being more sales to these customer segments but this is not a major driver for the Budget-Older families segment.

Higher sales may also be driven by more units of chips being bought per customer. Let's have a look at this next.

```
#Average number of units per customer by LIFESTAGE and
PREMIUM_CUSTOMER
avg_units_per_customer =merged_data.groupby(['LIFESTAGE',
'PREMIUM_CUSTOMER']).apply(
    lambda x: pd.Series({
        'AVG': x['PROD_QTY'].sum() / x['LYLTY_CARD_NBR'].nunique()
    })).reset_index()

# Plot average number of units per customer by LIFESTAGE and
PREMIUM_CUSTOMER
plt.figure(figsize=(8,4))
sns.barplot(data=avg_units_per_customer, x='LIFESTAGE', y='AVG',
hue='PREMIUM_CUSTOMER')
```

```
plt.title('Average number of units per customer by LIFESTAGE and PREMIUM_CUSTOMER')
plt.xticks(rotation=70)
plt.ylabel('Average Units')
plt.show()
```

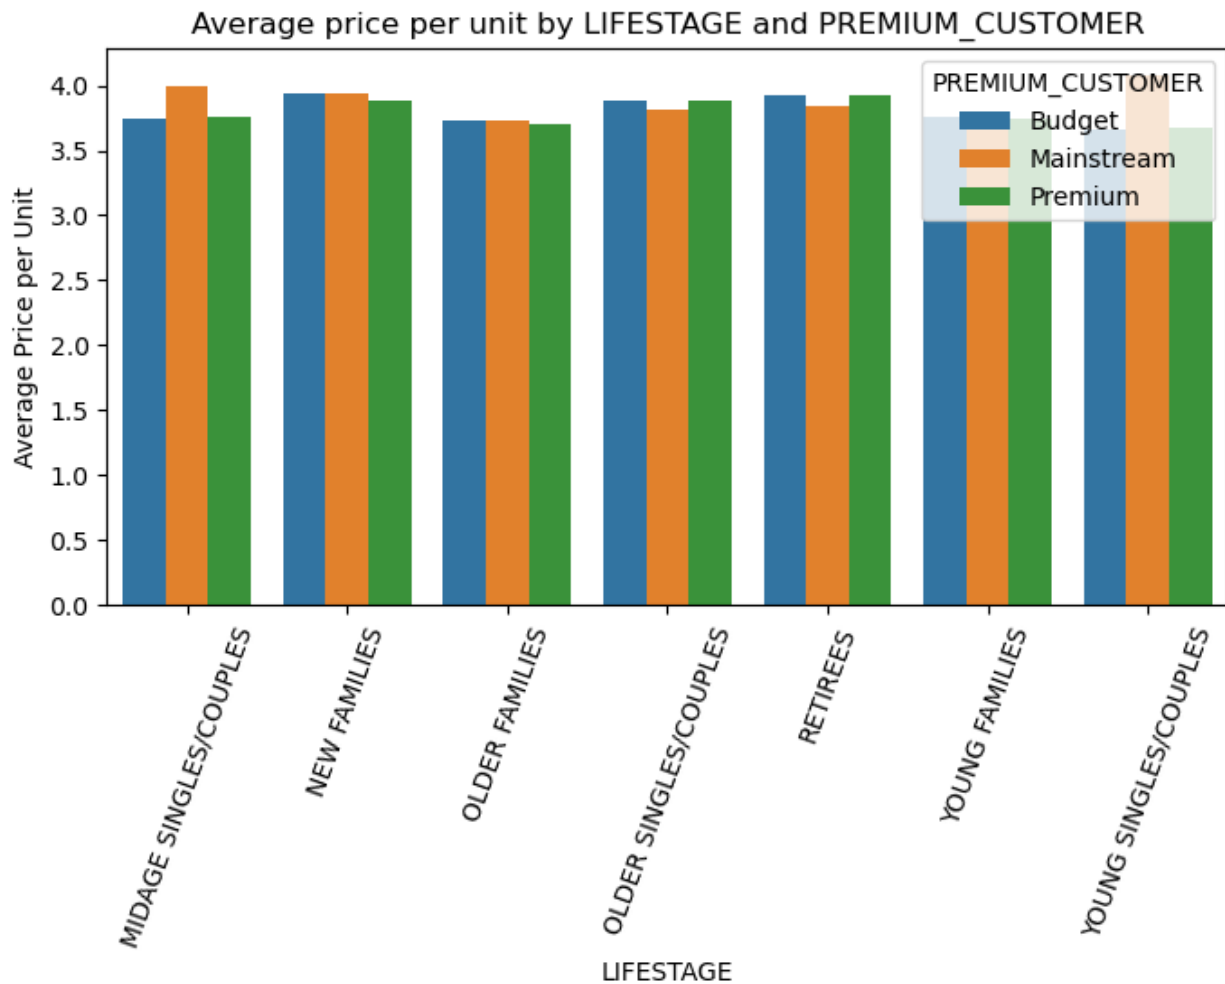


Older families and young families in general buy more chips per customer

Let's also investigate the average price per unit chips bought for each customer segment as this is also a driver of total sales.

```
# Average price per unit by LIFESTAGE and PREMIUM_CUSTOMER
avg_price_per_unit = merged_data.groupby(['LIFESTAGE',
    'PREMIUM_CUSTOMER']).apply(
    lambda x: pd.Series({
        'AVG': x['TOT_SALES'].sum() / x['PROD_QTY'].sum()
    })
).reset_index()
# Plot average price per unit by LIFESTAGE and PREMIUM_CUSTOMER
plt.figure(figsize=(8,4))
```

```
sns.barplot(data=avg_price_per_unit, x='LIFESTAGE', y='AVG',
hue='PREMIUM_CUSTOMER')
plt.title('Average price per unit by LIFESTAGE and PREMIUM_CUSTOMER')
plt.xticks(rotation=70)
plt.ylabel('Average Price per Unit')
plt.show()
```



Mainstream midage and young singles and couples are more willing to pay more per packet of chips compared to their budget and premium counterparts. This may be due to premium shoppers being more likely to buy healthy snacks and when they buy chips, this is mainly for entertainment purposes rather than their own consumption. This is also supported by there being fewer premium midage and young singles and couples buying chips compared to their mainstream counterparts.

As the difference in average price per unit isn't large, we can check if this difference is statistically different.

```
from scipy.stats import ttest_ind
```

```

# Calculate the price per unit
merged_data['price'] = merged_data['TOT_SALES'] /
merged_data['PROD_QTY']

# Filter data for the two groups
group_mainstream = merged_data[
    (merged_data['LIFESTAGE'].isin(['YOUNG SINGLES/COUPLES', 'MIDAGE
SINGLES/COUPLES'])) &
    (merged_data['PREMIUM_CUSTOMER'] == 'Mainstream')]
[['price']]

group_other = merged_data[
    (merged_data['LIFESTAGE'].isin(['YOUNG SINGLES/COUPLES', 'MIDAGE
SINGLES/COUPLES'])) &
    (merged_data['PREMIUM_CUSTOMER'] != 'Mainstream')]
[['price']]

# Perform independent t-test
t_stat, p_value = ttest_ind(group_mainstream, group_other,
alternative='greater')

# Display the result
print(f"T-statistic: {t_stat}, P-value: {p_value}")

T-statistic: 40.758905548288816, P-value: 0.0

```

The t-test results in a p-value is 0.000, i.e. the unit price for mainstream, young and mid-age singles and couples are significantly higher than that of budget or premium, young and midage singles and couples.

```

# Filter data into segment1 and other
segment1 = merged_data[(merged_data['LIFESTAGE'] == 'YOUNG
SINGLES/COUPLES') & (merged_data['PREMIUM_CUSTOMER'] == 'Mainstream')]
other = merged_data[~((merged_data['LIFESTAGE'] == 'YOUNG
SINGLES/COUPLES') & (merged_data['PREMIUM_CUSTOMER'] ==
'Mainstream'))]

# Calculate total quantity for each segment
quantity_segment1 = segment1['PROD_QTY'].sum()
quantity_other = other['PROD_QTY'].sum()

# Calculate the quantity per brand for each segment
quantity_segment1_by_brand = segment1.groupby('BRAND_NAME')
[['PROD_QTY']].sum().reset_index()
quantity_segment1_by_brand['targetSegment'] =
quantity_segment1_by_brand['PROD_QTY'] / quantity_segment1
quantity_segment1_by_brand = quantity_segment1_by_brand[['BRAND_NAME',
'targetSegment']]

quantity_other_by_brand = other.groupby('BRAND_NAME')

```

```

['PROD_QTY'].sum().reset_index()
quantity_other_by_brand['other'] = quantity_other_by_brand['PROD_QTY']
/ quantity_other
quantity_other_by_brand = quantity_other_by_brand[['BRAND_NAME',
'other']]

# Merge the two dataframes
brand_proportions = pd.merge(quantity_segment1_by_brand,
quantity_other_by_brand, on='BRAND_NAME')

# Calculate brand affinity
brand_proportions['affinityToBrand'] =
brand_proportions['targetSegment'] / brand_proportions['other']

# Sort by affinityToBrand in descending order
brand_proportions =
brand_proportions.sort_values(by='affinityToBrand', ascending=False)

# Display the result
print(brand_proportions)

```

	BRAND_NAME	targetSegment	other	affinityToBrand
19	TYRRELLS	0.029745	0.023961	1.241357
18	TWISTIES	0.043303	0.035151	1.231928
17	TOSTITOS	0.042678	0.035189	1.212818
9	KETTLE	0.184610	0.153860	1.199852
11	OLD	0.041481	0.034667	1.196575
12	PRINGLES	0.111848	0.093554	1.195550
5	DORITOS	0.122934	0.105205	1.168519
4	COBS	0.041872	0.036399	1.150350
8	INFUZIONI	0.060843	0.053252	1.142547
16	THINS	0.056783	0.053001	1.071353
7	GRNWVES	0.030603	0.029100	1.051658
3	CHEEZELS	0.016941	0.017321	0.978048
14	SMITHS	0.093606	0.121918	0.767781
6	FRENCH	0.003721	0.005400	0.689121
2	CHEETOS	0.007573	0.011200	0.676151
13	RRD	0.045359	0.068534	0.661842
10	NATURAL	0.018477	0.028763	0.642368
1	CCS	0.010539	0.017721	0.594742
15	SUNBITES	0.005985	0.011798	0.507313
20	WOOLWORTHS	0.028339	0.057818	0.490148
0	BURGER	0.002758	0.006186	0.445895

We can see that :

- Mainstream young singles/couples are 23% more likely to purchase Tyrrells chips compared to the rest of the population

- Mainstream young singles/couples are 56% less likely to purchase Burger Rings compared to the rest of the population Let's also find out if our target segment tends to buy larger packs of chips

Let's also find out if our target segment tends to buy larger packs of chips.

```
# Calculate the quantity per pack size for each segment
quantity_segment1_by_pack = segment1.groupby('PACK_SIZE')
['PROD_QTY'].sum().reset_index()
quantity_segment1_by_pack['targetSegment'] =
quantity_segment1_by_pack['PROD_QTY'] / quantity_segment1
quantity_segment1_by_pack = quantity_segment1_by_pack[['PACK_SIZE',
'targetSegment']]

quantity_other_by_pack = other.groupby('PACK_SIZE')
['PROD_QTY'].sum().reset_index()
quantity_other_by_pack['other'] = quantity_other_by_pack['PROD_QTY'] /
quantity_other
quantity_other_by_pack = quantity_other_by_pack[['PACK_SIZE',
'other']]

# Merge the two dataframes
pack_proportions = pd.merge(quantity_segment1_by_pack,
quantity_other_by_pack, on='PACK_SIZE')

# Calculate affinity to each pack size
pack_proportions['affinityToPack'] = pack_proportions['targetSegment']
/ pack_proportions['other']

# Sort by affinityToPack in descending order
pack_proportions = pack_proportions.sort_values(by='affinityToPack',
ascending=False)

# Display the result
print(pack_proportions)
```

	PACK_SIZE	targetSegment	other	affinityToPack
18	270	0.029875	0.023262	1.284299
21	380	0.030161	0.023756	1.269607
20	330	0.057538	0.046566	1.235605
15	210	0.014469	0.011779	1.228397
5	134	0.111848	0.093554	1.195550
3	110	0.099852	0.083721	1.192666
17	250	0.013428	0.011889	1.129460
6	135	0.013688	0.012222	1.120011
10	170	0.075728	0.075340	1.005144
11	175	0.236214	0.245645	0.961606
19	300	0.054909	0.057330	0.957778
7	150	0.150805	0.157764	0.955895
9	165	0.052333	0.058069	0.901226

0	0	0.019908	0.023154	0.859818
13	190	0.007052	0.011606	0.607648
12	180	0.003253	0.005668	0.573912
8	160	0.006037	0.011604	0.520297
2	90	0.005985	0.011798	0.507313
4	125	0.002837	0.005661	0.501024
14	200	0.008458	0.017496	0.483390
1	70	0.002863	0.005929	0.482781
16	220	0.002758	0.006186	0.445895

It looks like Mainstream young singles/couples are 27% more likely to purchase a 270g pack of chips compared to the rest of the population but let's dive into what brands sell this pack size.

```
# Filter data where PACK_SIZE is 270 and get unique product names
merged_data[merged_data['PACK_SIZE'] == 270]['PROD_NAME'].unique()

array(['Twisties Cheese 270g', 'Twisties Chicken270g'],
dtype=object)
```

Twisties are the only brand offering 270g packs and so this may instead be reflecting a higher likelihood of purchasing Twisties.

## Customer Lifetime Value (CLV)

```
# Calculate CLV for each customer
clv = merged_data.groupby('LYLTY_CARD_NBR').agg({
    'TOT_SALES': 'sum',
    'PROD_QTY': 'sum'
}).reset_index()
clv.columns = ['LYLTY_CARD_NBR', 'Total_Sales', 'Total_Quantity']

# Merge CLV with customer data
clv = clv.merge(purchase_behaviour, on='LYLTY_CARD_NBR', how='left')

# Group by customer segment to get average CLV
clv_segment = clv.groupby(['LIFESTAGE', 'PREMIUM_CUSTOMER']).agg({
    'Total_Sales': 'mean',
    'Total_Quantity': 'mean'
}).reset_index()

print(clv_segment)
```


		LIFESTAGE	PREMIUM_CUSTOMER	Total_Sales
Total_Quantity				
0	MIDAGE	SINGLES/COUPLES	Budget	23.493133
6.284667				
1	MIDAGE	SINGLES/COUPLES	Mainstream	27.002998
6.754796				
2	MIDAGE	SINGLES/COUPLES	Premium	23.875545
6.349239				



3	NEW FAMILIES	Budget	19.647347
4.994604			
4	NEW FAMILIES	Mainstream	19.937750
5.062426			
5	NEW FAMILIES	Premium	19.473639
5.013605			
6	OLDER FAMILIES	Budget	35.768305
9.585476			
7	OLDER FAMILIES	Mainstream	36.191416
9.718121			
8	OLDER FAMILIES	Premium	35.183627
9.503081			
9	OLDER SINGLES/COUPLES	Budget	27.554184
7.103168			
10	OLDER SINGLES/COUPLES	Mainstream	26.890130
7.061738			
11	OLDER SINGLES/COUPLES	Premium	27.633737
7.106737			
12	RETIREEES	Budget	25.227775
6.419326			
13	RETIREEES	Mainstream	23.895994
6.222016			
14	RETIREEES	Premium	25.045581
6.387080			
15	YOUNG FAMILIES	Budget	34.437129
9.178919			
16	YOUNG FAMILIES	Mainstream	33.749211
9.115921			
17	YOUNG FAMILIES	Premium	34.318552
9.155492			
18	YOUNG SINGLES/COUPLES	Budget	16.120906
4.397191			
19	YOUNG SINGLES/COUPLES	Mainstream	19.414829
4.760530			
20	YOUNG SINGLES/COUPLES	Premium	16.116485
4.386470			

*# Total Sales by LIFESTAGE and PREMIUM\_CUSTOMER*

```
plt.figure(figsize=(8, 4))
sns.barplot(data=clv_segment, x='LIFESTAGE', y='Total_Sales',
hue='PREMIUM_CUSTOMER')
plt.title('Total Sales by LIFESTAGE and PREMIUM_CUSTOMER')
plt.xticks(rotation=45)
plt.ylabel('Total Sales')
plt.show()
```



- Older Families and Young Families lead in total sales and quantities, suggesting they are the most valuable customer segments.
- Minimal differences in purchasing behavior across Budget, Mainstream, and Premium statuses within Older and Young Families.

## Churn Analysis

```
# Identify customers with last purchase date
last_purchase = merged_data.groupby('LYLTY_CARD_NBR')
['DATE'].max().reset_index()
last_purchase.columns = ['LYLTY_CARD_NBR', 'Last_Purchase_Date']

# Define churned customers (e.g., no purchase in the last 3 months of data)
end_date = merged_data['DATE'].max()
churn_threshold = end_date - pd.DateOffset(months=3)
churned_customers = last_purchase[last_purchase['Last_Purchase_Date']
< churn_threshold]

# Merge with customer data to analyze churned customers
churned_customers = churned_customers.merge(purchase_behaviour,
on='LYLTY_CARD_NBR', how='left')
churned_segment = churned_customers.groupby(['LIFESTAGE',
```

```
'PREMIUM_CUSTOMER'])).size().reset_index(name='Churn_Count')
```

```
print(churned_segment)
```

		LIFESTAGE	PREMIUM_CUSTOMER	Churn_Count
0	MIDAGE	SINGLES/COUPLES	Budget	669
1	MIDAGE	SINGLES/COUPLES	Mainstream	1422
2	MIDAGE	SINGLES/COUPLES	Premium	1102
3		NEW FAMILIES	Budget	532
4		NEW FAMILIES	Mainstream	411
5		NEW FAMILIES	Premium	306
6		OLDER FAMILIES	Budget	1601
7		OLDER FAMILIES	Mainstream	999
8		OLDER FAMILIES	Premium	791
9	OLDER	SINGLES/COUPLES	Budget	2016
10	OLDER	SINGLES/COUPLES	Mainstream	2011
11	OLDER	SINGLES/COUPLES	Premium	1938
12		RETIREEES	Budget	1948
13		RETIREEES	Mainstream	2889
14		RETIREEES	Premium	1687
15		YOUNG FAMILIES	Budget	1456
16		YOUNG FAMILIES	Mainstream	1024
17		YOUNG FAMILIES	Premium	833
18	YOUNG	SINGLES/COUPLES	Budget	2082
19	YOUNG	SINGLES/COUPLES	Mainstream	4276
20	YOUNG	SINGLES/COUPLES	Premium	1405

```
# Churn count by LIFESTAGE and PREMIUM_CUSTOMER
```

```
plt.figure(figsize=(8, 4))
```

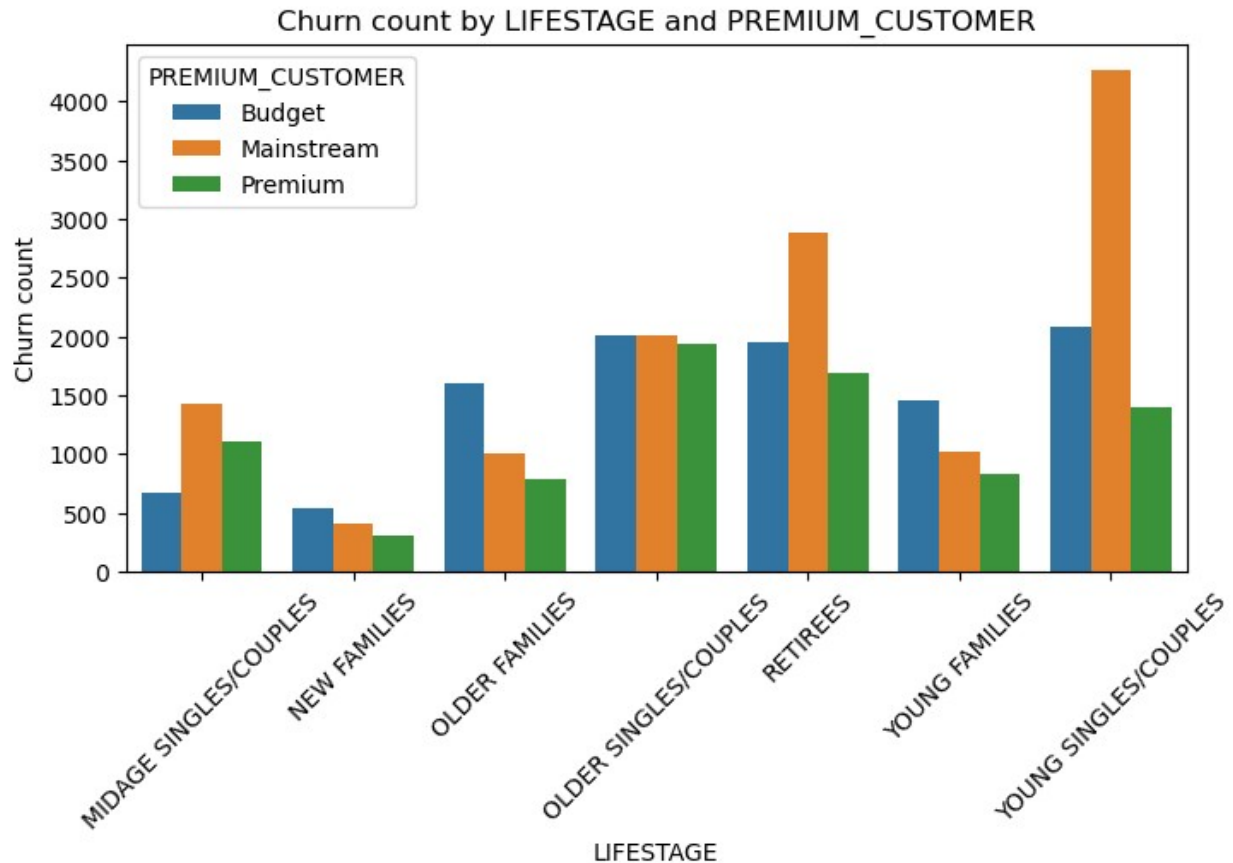
```
sns.barplot(data=churned_segment, x='LIFESTAGE', y='Churn_Count',  
hue='PREMIUM_CUSTOMER')
```

```
plt.title('Churn count by LIFESTAGE and PREMIUM_CUSTOMER')
```

```
plt.xticks(rotation=45)
```

```
plt.ylabel('Churn count')
```

```
plt.show()
```



- Young Singles/Couples have the highest churn rates, especially in the Mainstream category, indicating a need for targeted retention strategies.
- Retirees show substantial churn, with Mainstream customers leading, suggesting potential for improved engagement with senior customers.
- New and Young Families have comparatively lower churn rates, indicating better retention or satisfaction levels within these segments.

## Seasonality Analysis

```
# Ensure TOT_SALES is numeric
merged_data['TOT_SALES'] = pd.to_numeric(merged_data['TOT_SALES'],
errors='coerce')

# Extract Month from transaction date
merged_data['Month'] =
pd.to_datetime(merged_data['DATE']).dt.to_period('M')


# Group by Month and calculate total sales
monthly_sales = merged_data.groupby('Month').agg({'TOT_SALES':
'sum'}).reset_index()

# Convert Month back to a datetime object for plotting
```

```
monthly_sales['Month'] = monthly_sales['Month'].dt.to_timestamp()  
monthly_sales
```

	Month	TOT_SALES
0	2018-07-01	165258.20
1	2018-08-01	150688.30
2	2018-09-01	160522.00
3	2018-10-01	164415.70
4	2018-11-01	160233.70
5	2018-12-01	167913.40
6	2019-01-01	162642.30
7	2019-02-01	150645.50
8	2019-03-01	166265.20
9	2019-04-01	159845.10
10	2019-05-01	149923.95
11	2019-06-01	160538.60

```
# Plot monthly sales  
plt.figure(figsize=(8, 4))  
sns.lineplot(data=monthly_sales, x='Month', y='TOT_SALES')  
plt.title('Monthly Sales')  
plt.xlabel('Month')  
plt.ylabel('Total Sales')  
plt.show()
```



**Sales Fluctuation:** The sales figures fluctuate month-to-month, indicating variability in customer purchasing behavior across different times of the year.

**Peak Sales:** December 2018 shows the highest total sales (167,913.40), suggesting a potential seasonal peak during the holiday season.

**Low Sales Periods:** February 2019 and May 2019 have the lowest total sales (150,645.50 and 149,923.95 respectively), indicating potential post-holiday and pre-summer lulls.

**Moderate Variability:** While there are peaks and troughs, the overall variability in monthly sales is moderate, with values ranging from around 149,924 to 167,913.

**No Strong Seasonal Trend:** Aside from the holiday season peak, there is no strong, clear seasonal trend, indicating that other factors might also significantly influence monthly sales.

## Overall Summary:

The analysis conducted on the transaction data and purchase behavior data revealed several key insights and findings that can help drive business decisions and strategies. Here are the high-level findings and key callouts:

**Customer Segmentation:** Older Families and Young Families are the most valuable customer segments in terms of total sales and quantities. Mainstream young singles/couples are more likely to purchase Tyrrells chips, while they are less likely to purchase Burger Rings. Mainstream midage and young singles/couples are willing to pay more per packet of chips compared to their budget and premium counterparts.

**Churn Analysis:** Young Singles/Couples, especially in the Mainstream category, have the highest churn rates, indicating a need for targeted retention strategies. Retirees show substantial churn, with Mainstream customers leading, suggesting potential for improved engagement with senior customers. New and Young Families have comparatively lower churn rates, indicating better retention or satisfaction levels within these segments.

**Seasonality Analysis:** Sales figures fluctuate month-to-month, indicating variability in customer purchasing behavior. December 2018 shows the highest total sales, suggesting a potential seasonal peak during the holiday season. February 2019 and May 2019 have the lowest total sales, indicating potential post-holiday and pre-summer lulls.

**Customer Lifetime Value (CLV):** Older Families and Young Families lead in total sales and quantities, suggesting they are the most valuable customer segments. Minimal differences in purchasing behavior across Budget, Mainstream, and Premium statuses within Older and Young Families.

Overall, the analysis provides valuable insights into customer behavior, preferences, and trends, which can be leveraged to optimize marketing strategies, improve customer retention, and drive business growth. It is recommended to focus on targeted marketing efforts for high-churn segments and capitalize on seasonal trends to maximize sales opportunities.

